



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

---

Experiment No. 2
Analyze the Titanic Survival Dataset and apply appropriate regression technique
Date of Performance: 31—07—23
Date of Submission: 7—08—23



**Aim:** Analyze the Titanic Survival Dataset and apply appropriate Regression Technique.

**Objective:** Able to perform various feature engineering tasks, apply logistic regression on the given dataset and maximize the accuracy.

**Theory:**


Logistic Regression was used in the biological sciences in early twentieth century. It was then used in many social science applications. Logistic Regression is used when the dependent variable(target) is categorical and is binary in nature. In order to perform binary classification the logistic regression techniques makes use of Sigmoid function.

For example,

To predict whether an email is spam (1) or (0)

Whether the tumor is malignant (1) or not (0)

Consider a scenario where we need to classify whether an email is spam or not. If we use linear regression for this problem, there is a need for setting up a threshold based on which classification can be done. Say if the actual class is malignant, predicted continuous value 0.4 and the threshold value is 0.5, the data point will be classified as not malignant which can lead to serious consequence in real time.



From this example, it can be inferred that linear regression is not suitable for classification problem. Linear regression is unbounded, and this brings logistic regression into picture. Their value strictly ranges from 0 to 1.

#### **Dataset:**

The sinking of the Titanic is one of the most infamous shipwrecks in history.

On April 15, 1912, during her maiden voyage, the widely considered “unsinkable” RMS Titanic sank after colliding with an iceberg. Unfortunately, there weren’t enough lifeboats for everyone onboard, resulting in the death of 1502 out of 2224 passengers and crew.

While there was some element of luck involved in surviving, it seems some groups of people were more likely to survive than others.

In this challenge, we ask you to build a predictive model that answers the question: “what sorts of people were more likely to survive?” using passenger data (ie name, age, gender, socioeconomic class, etc).



# Vidyavardhini's College of Engineering & Technology

## Department of Computer Engineering

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

Variable Notes pclass: A proxy for socio-economic status (SES) 1st = Upper, 2nd

= Middle, 3rd = Lower age: Age is fractional if less than 1. If the age is estimated,

is it in the form of xx.5



# Vidyavardhini's College of Engineering & Technology

## Department of Computer Engineering

---

sibsp: The dataset defines family relations in this way...,

Sibling = brother, sister, stepbrother, stepsister

Spouse = husband, wife (mistresses and fiancés were ignored) parch:

The dataset defines family relations in this way..

Parent = mother, father

Child = daughter, son, stepdaughter, stepson

Some children travelled only with a nanny, therefore parch=0 for them



Vidyavardhini's College of Engineering & Technology  
Department of Computer Engineering

---

**Code:**



### **Conclusion:**

Passenger class (Pclass), gender (Sex), age (Age), number of siblings/spouses aboard (SibSp), and number of parents/children (Parch) are the features used for model construction. These characteristics are important because they may have an impact on survival rates and represent socioeconomic factors, such as giving higher-class passengers priority, prioritising women during evacuations, giving children and the elderly priority, allowing family members to travel, and relationships between departure port and socioeconomic backgrounds. An accuracy score of 0.8076 for the training data shows that the model correctly forecasts survival outcomes for this dataset. According to the test data accuracy score of 0.7821, the model performs well with fresh, untested data. These accuracy values indicate that the model successfully generalizes novel data.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import warnings
```

```
data = pd.read_csv("/content/train (2).csv")
```

```
print(data)
```

↗

	PassengerId	Survived	Pclass	\								
0	1	0	3									
1	2	1	1									
2	3	1	3									
3	4	1	1									
4	5	0	3									
..	...	...	...									
886	887	0	2									
887	888	1	1									
888	889	0	3									
889	890	1	1									
890	891	0	3									

	Name	Sex	Age	SibSp	\							
0	Braund, Mr. Owen Harris	male	22.0	1								
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1								
2	Heikkinen, Miss. Laina	female	26.0	0								
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1								
4	Allen, Mr. William Henry	male	35.0	0	..							
	...	...	...	...								
886	Montvila, Rev. Juozas	male	27.0	0								
887	Graham, Miss. Margaret Edith	female	19.0	0								
888	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1								
889	Behr, Mr. Karl Howell	male	26.0	0								
890	Dooley, Mr. Patrick	male	32.0	0								

	Parch	Ticket	Fare	Cabin	Embarked							
0	0	A/5 21171	7.2500	NaN	S							
1	0	PC 17599	71.2833	C85	C							
2	0	STON/O2. 3101282	7.9250	NaN	S							
3	0	113803	53.1000	C123	S							
4	0	373450	8.0500	NaN	S	..	...					
886	0	211536	13.0000	NaN	S							
887	0	112053	30.0000	B42	S							
888	2	W./C. 6607	23.4500	NaN	S							
889	0	111369	30.0000	C148	C							
890	0	370376	7.7500	NaN	Q							

[891 rows x 12 columns]

```
data.shape
```

(891, 12)

```
data.info() # getting some informations about the data
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null   int64
1   Survived    891 non-null   int64
2   Pclass      891 non-null   int64
3   Name        891 non-null   object
4   Sex         891 non-null   object
5   Age         714 non-null   float64
6   SibSp       891 non-null   int64
7   Parch       891 non-null   int64
8   Ticket      891 non-null   object
9   Fare        891 non-null   float64
10  Cabin       204 non-null   object
11  Embarked    889 non-null   object
```



```
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
data.isnull().sum() # check the number of missing values in each column
```

```
PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age          177
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin        687
Embarked       2
dtype: int64
```

```
data = data.drop(columns='Cabin', axis=1)
```

```
data['Age'].fillna(data['Age'].mean(), inplace=True) # replacing the missing values in "Age" column with mean value
```

```
print(data['Embarked'].mode()) # finding the mode value of "Embarked" column
```

```
0    S
Name: Embarked, dtype: object
```

```
print(data['Embarked'].mode()[0])
```

```
S
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True) # replacing the missing values in "Embarked" column with mode value
```

```
data.isnull().sum() # check the number of missing values in each column
```

```
PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
dtype: int64
```

```
data.describe()
```


	PassengerId	Survived	Pclass	Age	SibSp	Parch	
<b>count</b>	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000	891.0
<b>mean</b>	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.2
<b>std</b>	257.353842	0.486592	0.836071	13.002015	1.102743	0.806057	49.6
<b>min</b>	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.0
<b>25%</b>	223.500000	0.000000	2.000000	22.000000	0.000000	0.000000	7.9
<b>50%</b>	446.000000	0.000000	3.000000	29.699118	0.000000	0.000000	14.4
<b>75%</b>	668.500000	1.000000	3.000000	35.000000	1.000000	0.000000	31.0
<b>max</b>	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.3

```
data['Survived'].value_counts() # finding the number of people survived and not survived
```

```
0    549
1    342
Name: Survived, dtype: int64
data['Sex'].value_counts()

male      577 female
314 Name: Sex, dtype:
int64
```

```
# number of survivors Gender wise
# 1st male and other female # 0 are the one who
did not survived sns.countplot(x='Sex',
hue='Survived', data=data)
```



```
data['Embarked'].value_counts()

S    646
C    168
Q     77
Name: Embarked, dtype: int64
```

```
# converting categorical Columns
data.replace({'Sex':{'male':0,'female':1}, 'Embarked':{'S':0,'C':1,'Q':2}}, inplace=True)
```

```
X = data.drop(columns = ['PassengerId','Name','Ticket','Survived'],axis=1)
Y = data['Survived']
```

```
print(X)
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	0	22.000000	1	0	7.2500	0
1	1	1	38.000000	1	0	71.2833	1
2	3	1	26.000000	0	0	7.9250	0
3	1	1	35.000000	1	0	53.1000	0
4	3	0	35.000000	0	0	8.0500	0
886	2	0	27.000000	0	0	13.0000	0
887	1	1	19.000000	0	0	30.0000	0
888	3	1	29.699118	1	2	23.4500	0
889	1	0	26.000000	0	0	30.0000	1
890	3	0	32.000000	0	0	7.7500	2
891	rows x 7 columns						

```
print(Y)
```

0	0
1	1
2	1
3	1

```
4      0      ..
886    0
887    1
888    0
889    1
890    0
Name: Survived, Length: 891, dtype: int64

data.head()

PassengerId Survived Pclass      Name Sex Age SibSp Parch      Ticket      F
-----
0           1         0      3  Braund, Mr. Owen Harris 0      22.0      1      0      A/5 21171 7.2
1           2         1      1  Cumings, Mrs. John Bradley (Florence Briggs) 1      38.0      1      0      PC 17599 71.2

X.head()

Pclass Sex Age SibSp Parch      Fare Embarked
-----
0       3   0  22.0      1      0      7.2500  0
1       1   1  38.0      1      0     71.2833  1
2       3   1  26.0      0      0      7.9250  0
3       1   1  35.0      1      0     53.1000  0
4       3   0  35.0      0      0      8.0500  0

Y.head()

0      0
1      1
2      1
3      1
4      0
Name: Survived, dtype: int64

#Splitting the data into training data & Test data
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2, random_state=2)
print(X.shape, X_train.shape, X_test.shape)

(891, 7) (712, 7) (179, 7)

logr = LogisticRegression()

# training the Logistic Regression model with training data
logr.fit(X_train, Y_train) /usr/local/lib/python3.10/dist-
```

packages/sklearn/linear\_model/\_logistic.py:458: Con STOP: TOTAL NO. of  
ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

# accuracy on training data [https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)<sup>a</sup> n

X\_train\_prediction = logit.predict(X\_train) n\_iter\_i = \_check\_optimize\_result(

training\_data\_accuracy = accuracy\_score(Y\_train, X\_train\_prediction) print("Accuracy score of training data:",

training\_data\_accuracy) LogisticRegression()

Accuracy score of training data: 0.8075842696629213

# accuracy on test data X\_test\_prediction =

logit.predict(X\_test) test\_data\_accuracy =

accuracy\_score(Y\_test, X\_test\_prediction) print('Accuracy

score of test data : ', test\_data\_accuracy)

Accuracy score of test data : 0.7821229050279329